

In [316]: `print("test")`

test

In [317]: `# First adding all necessary Libraries:`  
`import numpy as np`  
`import pandas as pd`  
`import seaborn as sns`  
`import matplotlib.pyplot as plt`  
`from sklearn.preprocessing import LabelEncoder, StandardScaler`  
`from sklearn.linear_model import LinearRegression, Lasso`  
`from sklearn.metrics import mean_squared_error, mean_absolute_error`  
`from sklearn.ensemble import RandomForestRegressor`  
`import warnings`  
`warnings.filterwarnings("ignore")`

In [318]: `# Loading the data of "LaptopPrice" dataset`  
`df = pd.read_csv('D:laptopPrice.csv')`

In [319]: `# display the first five records`  
`df.head(5)`

Out[319]:

	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	
0	ASUS	Intel	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	W
1	Lenovo	Intel	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	W
2	Lenovo	Intel	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	W
3	ASUS	Intel	Core i5	10th	8 GB	DDR4	512 GB	0 GB	W
4	ASUS	Intel	Celeron Dual	Not Available	4 GB	DDR4	0 GB	512 GB	W

In [320]: `#The shape function is used to display the total number of rows and columns of`  
`print(df.shape)`

(823, 19)

```
In [321]: #Checking for null values in each column and displaying the sum of all null values  
missing_values = df.isnull().sum()  
print("Missing Values:")  
print(missing_values)
```

```
Missing Values:  
brand                0  
processor_brand      0  
processor_name       0  
processor_gnrtn      0  
ram_gb               0  
ram_type             0  
ssd                  0  
hdd                  0  
os                   0  
os_bit               0  
graphic_card_gb     0  
weight              0  
warranty             0  
Touchscreen         0  
msoffice             0  
Price               0  
rating              0  
Number of Ratings   0  
Number of Reviews   0  
dtype: int64
```

```
In [322]: #Removing the rows with empty values
print(df.dropna())
```

	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	\
0	ASUS	Intel	Core i3	10th	4 GB	DDR4	
1	Lenovo	Intel	Core i3	10th	4 GB	DDR4	
2	Lenovo	Intel	Core i3	10th	4 GB	DDR4	
3	ASUS	Intel	Core i5	10th	8 GB	DDR4	
4	ASUS	Intel	Celeron Dual	Not Available	4 GB	DDR4	
..	...	...	...	...	...	...	
818	ASUS	AMD	Ryzen 9	Not Available	4 GB	DDR4	
819	ASUS	AMD	Ryzen 9	Not Available	4 GB	DDR4	
820	ASUS	AMD	Ryzen 9	Not Available	4 GB	DDR4	
821	ASUS	AMD	Ryzen 9	Not Available	4 GB	DDR4	
822	Lenovo	AMD	Ryzen 5	10th	8 GB	DDR4	

	ssd	hdd	os	os_bit	graphic_card_gb	weight	\
0	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	
1	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	
2	0 GB	1024 GB	Windows	64-bit	0 GB	Casual	
3	512 GB	0 GB	Windows	32-bit	2 GB	Casual	
4	0 GB	512 GB	Windows	64-bit	0 GB	Casual	
..	...	...	...	...	...	...	
818	1024 GB	0 GB	Windows	64-bit	0 GB	Casual	
819	1024 GB	0 GB	Windows	64-bit	0 GB	Casual	
820	1024 GB	0 GB	Windows	64-bit	4 GB	Casual	
821	1024 GB	0 GB	Windows	64-bit	4 GB	Casual	
822	512 GB	0 GB	DOS	64-bit	0 GB	ThinNlight	

	warranty	Touchscreen	msoffice	Price	rating	Number of Ratings	\
0	No warranty	No	No	34649	2 stars	3	
1	No warranty	No	No	38999	3 stars	65	
2	No warranty	No	No	39999	3 stars	8	
3	No warranty	No	No	69990	3 stars	0	
4	No warranty	No	No	26990	3 stars	0	
..	...	...	...	...	...	...	
818	1 year	No	No	135990	3 stars	0	
819	1 year	No	No	144990	3 stars	0	
820	1 year	No	No	149990	3 stars	0	
821	1 year	No	No	142990	3 stars	0	
822	No warranty	No	No	57490	4 stars	18	

	Number of Reviews
0	0
1	5
2	1
3	0
4	0
..	...
818	0
819	0
820	0
821	0
822	4

[823 rows x 19 columns]

```
In [323]: # Checking if there is any duplicates value
df.duplicated().sum()
```

Out[323]: 21

```
In [324]: df = df.drop_duplicates()
```

```
In [325]: # Display basic information about the dataset
# info() is a method used to provide a summary of dataframe
# and understand the dataset .Also getting the structure of dataframe that am g
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 802 entries, 0 to 822
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   brand                  802 non-null    object
1   processor_brand        802 non-null    object
2   processor_name         802 non-null    object
3   processor_gnrtn        802 non-null    object
4   ram_gb                 802 non-null    object
5   ram_type               802 non-null    object
6   ssd                    802 non-null    object
7   hdd                    802 non-null    object
8   os                     802 non-null    object
9   os_bit                 802 non-null    object
10  graphic_card_gb        802 non-null    object
11  weight                 802 non-null    object
12  warranty               802 non-null    object
13  Touchscreen            802 non-null    object
14  msoffice                802 non-null    object
15  Price                  802 non-null    int64
16  rating                 802 non-null    object
17  Number of Ratings      802 non-null    int64
18  Number of Reviews      802 non-null    int64
dtypes: int64(3), object(16)
memory usage: 125.3+ KB
```

```
In [326]: #Checking the data types to see if all the data is in correct format.  
# dtypes used to check and understanding the types of data presented in each column  
df.dtypes
```

```
Out[326]: brand                object  
processor_brand              object  
processor_name               object  
processor_gnrtn              object  
ram_gb                       object  
ram_type                     object  
ssd                          object  
hdd                          object  
os                           object  
os_bit                       object  
graphic_card_gb             object  
weight                       object  
warranty                     object  
Touchscreen                  object  
msoffice                     object  
Price                        int64  
rating                       object  
Number of Ratings            int64  
Number of Reviews            int64  
dtype: object
```

```
In [327]: # count the number of unique values in each column of a DataFrame.  
df.nunique()
```

```
Out[327]: brand                8  
processor_brand              3  
processor_name               11  
processor_gnrtn              8  
ram_gb                       4  
ram_type                     6  
ssd                          7  
hdd                          4  
os                           3  
os_bit                       2  
graphic_card_gb             5  
weight                       3  
warranty                     4  
Touchscreen                  2  
msoffice                     2  
Price                        405  
rating                       5  
Number of Ratings            282  
Number of Reviews            135  
dtype: int64
```

```
In [328]: #For numerical columns only
df.describe()
```

Out[328]:

	Price	Number of Ratings	Number of Reviews
count	802.000000	802.00000	802.000000
mean	76625.543641	299.84414	36.089776
std	45232.984422	1001.78442	118.313553
min	16990.000000	0.00000	0.000000
25%	45990.000000	0.00000	0.000000
50%	63990.000000	17.00000	2.000000
75%	89525.000000	140.25000	18.000000
max	441990.000000	15279.00000	1947.000000

```
In [329]: df.duplicated().sum()
```

Out[329]: 0

```
In [330]: # Checking the number of numeric features and cat_features (Categorical) from c
numeric_features = [feature for feature in df.columns if df[feature].dtype !=
cat_features = [feature for feature in df.columns if df[feature].dtype == 'obje
# Display Numerical and Categorical variables
print(" Numerical features: ", numeric_features)
print("Categorical features:", cat_features)
```

Numerical features: ['Price', 'Number of Ratings', 'Number of Reviews']  
Categorical features: ['brand', 'processor\_brand', 'processor\_name', 'processo  
r\_gnrtn', 'ram\_gb', 'ram\_type', 'ssd', 'hdd', 'os', 'os\_bit', 'graphic\_card\_g  
b', 'weight', 'warranty', 'Touchscreen', 'msoffice', 'rating']

```
In [331]: # describe () method used to describe numerical column only
df.describe(include = 'object')
```

Out[331]:

	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd
count	802	802	802	802	802	802	802	802
unique	8	3	11	8	4	6	7	4
top	ASUS	Intel	Core i5	11th	8 GB	DDR4	512 GB	0 GB
freq	243	594	284	328	404	690	389	602

```
In [332]: # Looking for minimum Number of Ratings
df.loc[df['Number of Ratings'] == 1, 'Number of Ratings'] = 10
```

## df.describe()

In [333]: `df['Price'].describe()`

```
Out[333]: count      802.000000
mean      76625.543641
std       45232.984422
min       16990.000000
25%       45990.000000
50%       63990.000000
75%       89525.000000
max       441990.000000
Name: Price, dtype: float64
```

In [334]: `df.describe(include = 'object')` *#summary statistics for categorical values*

```
Out[334]:
```


	brand	processor_brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd
<b>count</b>	802	802	802	802	802	802	802	802
<b>unique</b>	8	3	11	8	4	6	7	4
<b>top</b>	ASUS	Intel	Core i5	11th	8 GB	DDR4	512 GB	0 GB
<b>freq</b>	243	594	284	328	404	690	389	602

```
In [335]: numeric_features = [feature for feature in df.columns if df[feature].dtype !=
cat_features = [feature for feature in df.columns if df[feature].dtype == 'object']
print("Numerical features: ", numeric_features)
print("Categorical features:", cat_features)
```

```
Numerical features: ['Price', 'Number of Ratings', 'Number of Reviews']
Categorical features: ['brand', 'processor_brand', 'processor_name', 'processor_gnrtn', 'ram_gb', 'ram_type', 'ssd', 'hdd', 'os', 'os_bit', 'graphic_card_gb', 'weight', 'warranty', 'Touchscreen', 'msoffice', 'rating']
```

```
In [336]: import seaborn as sns
plt.figure(figsize=(10,6))
sns.regplot(x="Number of Ratings", y="Price", data=df)
```

Out[336]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d8859cd0>




```
In [337]: from scipy import stats
pearson_coef, p_value = stats.pearsonr(df['Number of Ratings'], df['Price'])
print("The Pearson Correlation Coefficient is", pearson_coef, " with a P-value
```

The Pearson Correlation Coefficient is -0.15246729283878194 with a P-value of P = 1.4479830870223598e-05




```
In [338]: plt.figure(figsize=(10,6))  
sns.regplot(x="Number of Reviews", y="Price", data=df)
```

Out[338]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d884ffd0>




```
In [339]: # In the given plot below, it is observed that the price range vary for ASUS ar  
# This indicates the categories can vary with price hence features can be used  
sns.boxplot(x="brand", y="Price", data=df)
```

Out[339]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d98f0820>



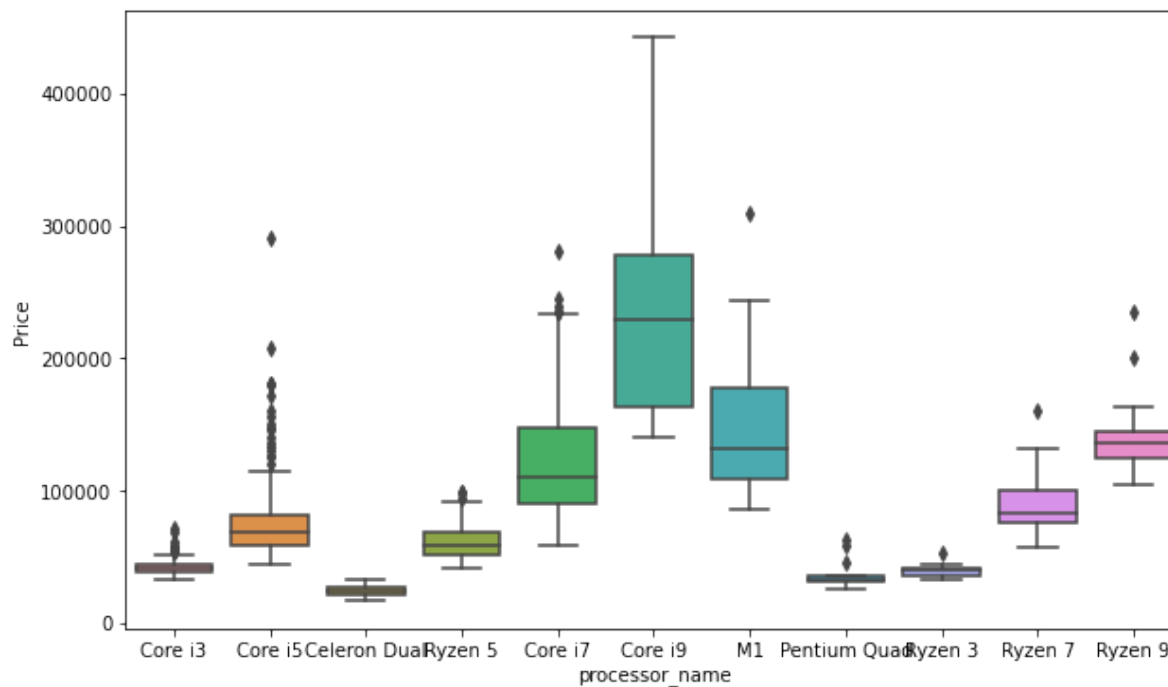
```
In [340]: plt.figure(figsize=(10,6))  
sns.boxplot(x="processor_brand", y="Price", data=df)
```

```
Out[340]: <matplotlib.axes._subplots.AxesSubplot at 0x122d98f0a90>
```




```
In [341]: plt.figure(figsize=(10,6))  
sns.boxplot(x="processor_name", y="Price", data=df)
```

```
Out[341]: <matplotlib.axes._subplots.AxesSubplot at 0x122d9a6d880>
```



```
In [342]: plt.figure(figsize=(10,6))  
sns.boxplot(x="processor_gnrtn", y="Price", data=df)
```


Out[342]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d98e7850>



In [343]: *# processor\_name feature shows a huge difference in price ranges between laptops*  
*# This feature is very important for price prediction as the bigger the difference, the more accurate the prediction*


```
In [344]: plt.figure(figsize=(10,6))  
sns.boxplot(x="ram_gb", y="Price", data=df)
```

Out[344]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d9a7f0d0>




```
In [345]: plt.figure(figsize=(10,6))  
sns.boxplot(x="ram_type", y="Price", data=df)
```

Out[345]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d9e02eb0>




```
In [346]: plt.figure(figsize=(10,6))  
sns.boxplot(x="ssd", y="Price", data=df)
```

Out[346]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d9ee97f0>




```
In [347]: plt.figure(figsize=(10,6))  
sns.boxplot(x="hdd", y="Price", data=df)
```

Out[347]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d9fcd730>




```
In [348]: plt.figure(figsize=(10,6))  
sns.boxplot(x="os", y="Price", data=df)
```

Out[348]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d9b99cd0>




```
In [349]: plt.figure(figsize=(10,6))  
sns.boxplot(x="os_bit", y="Price", data=df)
```

Out[349]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d7452dc0>




```
In [350]: plt.figure(figsize=(10,6))  
sns.boxplot(x="graphic_card_gb", y="Price", data=df)
```

Out[350]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d73e9bb0>




```
In [351]: plt.figure(figsize=(10,6))  
sns.boxplot(x="weight", y="Price", data=df)
```

Out[351]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d2941340>




```
In [352]: plt.figure(figsize=(10,6))  
sns.boxplot(x="warranty", y="Price", data=df)
```

Out[352]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122d9cb0130>




```
In [353]: plt.figure(figsize=(10,6))  
sns.boxplot(x="Touchscreen", y="Price", data=df)
```

Out[353]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122da027370>



```
In [354]: plt.figure(figsize=(10,6))  
sns.boxplot(x="msoffice", y="Price", data=df)
```


Out[354]: <matplotlib.axes.\_subplots.AxesSubplot at 0x122da095580>






```
In [355]: plt.figure(figsize=(10,6))  
sns.boxplot(x="os_bit", y="Price", data=df)
```

```
Out[355]: <matplotlib.axes._subplots.AxesSubplot at 0x122d5a28a00>
```



```
In [356]: plt.figure(figsize=(10,6))  
sns.boxplot(x="rating", y="Price", data=df)
```

```
Out[356]: <matplotlib.axes._subplots.AxesSubplot at 0x122da164760>
```



```
In [357]: df.drop(['weight', 'warranty', 'Touchscreen', 'processor_brand', 'os_bit'], axis
```

```
In [358]: df
```

Out[358]:

	brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	graphic
0	ASUS	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	
1	Lenovo	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	
2	Lenovo	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	
3	ASUS	Core i5	10th	8 GB	DDR4	512 GB	0 GB	Windows	
4	ASUS	Celeron Dual	Not Available	4 GB	DDR4	0 GB	512 GB	Windows	
...	...	...	...	...	...	...	...	...	...
818	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	
819	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	
820	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	
821	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	
822	Lenovo	Ryzen 5	10th	8 GB	DDR4	512 GB	0 GB	DOS	

802 rows x 14 columns



In [359]: df

Out[359]:

	brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	graphic
0	ASUS	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	
1	Lenovo	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	
2	Lenovo	Core i3	10th	4 GB	DDR4	0 GB	1024 GB	Windows	
3	ASUS	Core i5	10th	8 GB	DDR4	512 GB	0 GB	Windows	
4	ASUS	Celeron Dual	Not Available	4 GB	DDR4	0 GB	512 GB	Windows	
...	...	...	...	...	...	...	...	...	...
818	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	
819	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	
820	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	
821	ASUS	Ryzen 9	Not Available	4 GB	DDR4	1024 GB	0 GB	Windows	
822	Lenovo	Ryzen 5	10th	8 GB	DDR4	512 GB	0 GB	DOS	

802 rows × 14 columns

```

In [360]: #brand', 'processor_name', 'processor_gnrtn', 'ram_gb', 'ram_type', 'ssd', 'hdd', 'os', 'graphic_card_gb', 'msoffice', 'rating'
from sklearn.preprocessing import LabelEncoder

labelencoder = LabelEncoder()
df.brand = labelencoder.fit_transform(df.brand)
df.processor_name = labelencoder.fit_transform(df.processor_name)
df.processor_gnrtn = labelencoder.fit_transform(df.processor_gnrtn)
df.ram_gb = labelencoder.fit_transform(df.ram_gb)
df.ram_type = labelencoder.fit_transform(df.ram_type)
df.ssd = labelencoder.fit_transform(df.ssd)
df.hdd = labelencoder.fit_transform(df.hdd)
df.os = labelencoder.fit_transform(df.os)
df.graphic_card_gb = labelencoder.fit_transform(df.graphic_card_gb)
df.msoffice = labelencoder.fit_transform(df.msoffice)
df.rating = labelencoder.fit_transform(df.rating)
from sklearn.preprocessing import LabelEncoder

```

```
In [361]: import scipy.stats as stats
df = stats.zscore(df)
```

```
In [362]: df
```

```
Out[362]: array([[ -1.15409599, -0.87301913, -0.91832804, ..., -2.76166878,
        -0.29656123, -0.30522536],
       [ 0.87903961, -0.87301913, -0.91832804, ..., -1.00105005,
        -0.23463202, -0.2629384 ],
       [ 0.87903961, -0.87301913, -0.91832804, ..., -1.00105005,
        -0.29156694, -0.29676797],
       ...,
       [ -1.15409599,  2.20686947,  1.48628652, ..., -1.00105005,
        -0.2995578 , -0.30522536],
       [ -1.15409599,  2.20686947,  1.48628652, ..., -1.00105005,
        -0.2995578 , -0.30522536],
       [ 0.87903961,  1.52244978, -0.91832804, ...,  0.75956868,
        -0.28157836, -0.27139579]])
```

```
In [363]: x_train=df.iloc[:,0:13]
y_train=df.iloc[:,10]
```

```
-----
AttributeError                                Traceback (most recent call last)
<ipython-input-363-e365df6435b8> in <module>
----> 1 x_train=df.iloc[:,0:13]
      2 y_train=df.iloc[:,10]

AttributeError: 'numpy.ndarray' object has no attribute 'iloc'
```

```
In [364]: x_train.head()
```

```
Out[364]:
```

	brand	processor_name	processor_gnrtn	ram_gb	ram_type	ssd	hdd	os	graphic_card_gb
0	1	1	0	2	1	0	1	2	0
1	5	1	0	2	1	0	1	2	0
2	5	1	0	2	1	0	1	2	0
3	1	2	0	3	1	6	0	2	1
4	1	0	7	2	1	0	3	2	0

```
In [365]: y_train.head()
```

```
Out[365]: 0    34649
1    38999
2    39999
3    69990
4    26990
Name: Price, dtype: int64
```

```
In [366]: # importing train_test_split from sklearn
from sklearn.model_selection import train_test_split
# splitting the data # 30% for testing is used
X_train, X_test, Y_train, Y_test = train_test_split(x_train, y_train, test_size=
```

```
In [367]: #Multiple Linear Regression
from sklearn.linear_model import LinearRegression

model = LinearRegression()
model_mlr = model.fit(X_train,Y_train)
```

```
In [368]: #Making price prediction using the testing set (Fit to MLR)
Y_pred_MLR = model_mlr.predict(X_test)
```

```
In [369]: #Calculating the Mean Square Error for MLR model
mse_MLR = mean_squared_error(Y_test, Y_pred_MLR)
print('The mean square error for Multiple Linear Regression: ', mse_MLR)

The mean square error for Multiple Linear Regression:  3.4493088314991315e-22
```

```
In [370]: #The mean square error for Multiple Linear Regression:  0.3674647167443785
```

```
In [371]: #Calculating the Mean Absolute Error for MLR model
mae_MLR= mean_absolute_error(Y_test, Y_pred_MLR)
print('The mean absolute error for Multiple Linear Regression: ', mae_MLR)

The mean absolute error for Multiple Linear Regression:  1.5533112831939305e-11
```

```
In [372]: #Calling the random forest model and fitting the training data
rfModel = RandomForestRegressor()
model_rf = rfModel.fit(X_train,Y_train)
```

```
In [373]: #Prediction of Laptop prices using the testing data
Y_pred_RF = model_rf.predict(X_test)
```

```
In [374]: #Calculating the Mean Square Error for Random Forest Model
mse_RF = mean_squared_error(Y_test, Y_pred_RF)
print('The mean square error of price and predicted value is: ', mse_RF)

The mean square error of price and predicted value is:  5353991.133844395
```

```
In [375]: #Calculating the Mean Absolute Error for Random Forest Model
mae_RF= mean_absolute_error(Y_test, Y_pred_RF)
print('The mean absolute error of price and predicted value is: ', mae_RF)

The mean absolute error of price and predicted value is:  343.75962655601637
```

```
In [376]: #LASSO Model
#Calling the model and fitting the training data
LassoModel = Lasso()
model_lm = LassoModel.fit(X_train,Y_train)
```

```
In [377]: #Price prediction using testing data
Y_pred_lasso = model_lm.predict(X_test)
```

```
In [378]: #Mean Absolute Error for LASSO Model
mae_lasso= mean_absolute_error(Y_test, Y_pred_lasso)
print('The mean absolute error of price and predicted value is: ', mae_lasso)
```

The mean absolute error of price and predicted value is: 1.3399300380673692e-05

```
In [379]: #Mean Squared Error for the LASSO Model
mse_lasso = mean_squared_error(Y_test, Y_pred_lasso)
print('The mean square error of price and predicted value is: ', mse_lasso)
```

The mean square error of price and predicted value is: 3.446638548304528e-10

```
In [380]: scores = [('MLR', mae_MLR),
                    ('Random Forest', mae_RF),
                    ('LASSO', mae_lasso)
                  ]
```

```
In [381]: mae = pd.DataFrame(data = scores, columns=['Model', 'MAE Score'])
mae
```

Out[381]:

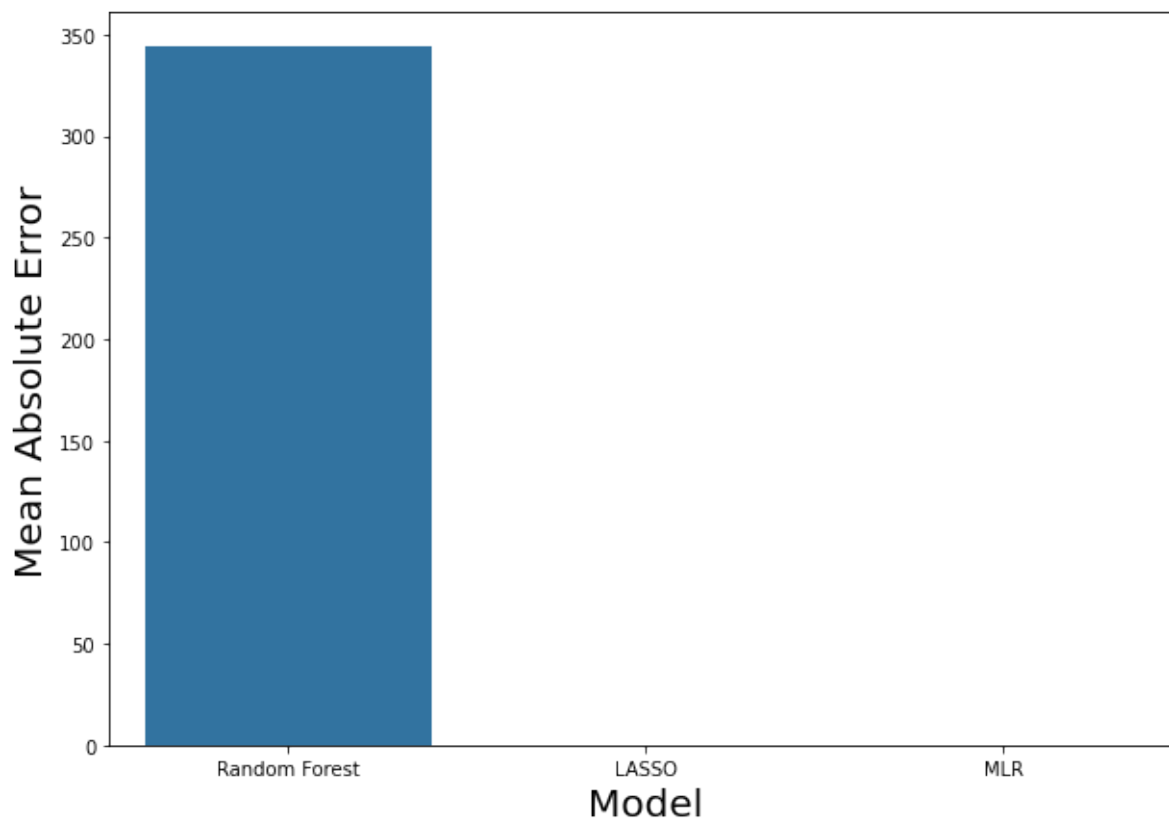
	Model	MAE Score
0	MLR	1.553311e-11
1	Random Forest	3.437596e+02
2	LASSO	1.339930e-05

```
In [382]: # By observing MAE score of MLR, Random Forest and LASSO I observe that random
# in general: the lower MAE score the better model.
# Hence , I conclude that Random Forest is the better model among the selected
```

```
In [383]: mae.sort_values(by=['MAE Score'], ascending=False, inplace=True)

f, axe = plt.subplots(1,1, figsize=(10,7))
sns.barplot(x = mae['Model'], y=mae['MAE Score'], ax = axe)
axe.set_xlabel('Model', size=20)
axe.set_ylabel('Mean Absolute Error', size=20)

plt.show()
```



In [ ]:

In [ ]: